

```
In [1]:
```

```
import pandas as pd
import numpy as np
from sklearn import preprocessing
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]:
```

```
data = pd.read_csv("Churn.csv",na_values='NaN')
data
```

```
Out[2]:
```

	Unnamed: 0	state	area.code	account.length	voice.plan	voice.messages	intl.plan	intl.mins
0	1	KS	area_code_415	128	yes	25	no	10.0
1	2	OH	area_code_415	107	yes	26	no	13.7
2	3	NJ	area_code_415	137	no	0	no	12.2
3	4	OH	area_code_408	84	no	0	yes	6.6
4	5	OK	area_code_415	75	no	0	yes	10.1
...
4995	4996	HI	area_code_408	50	yes	40	no	9.9
4996	4997	WV	area_code_415	152	no	0	no	14.7
4997	4998	DC	area_code_415	61	no	0	no	13.6
4998	4999	DC	area_code_510	109	no	0	no	8.5
4999	5000	VT	area_code_415	86	yes	34	no	9.3

5000 rows × 21 columns

```
In [3]:
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Unnamed: 0        5000 non-null   int64  
 1   state            5000 non-null   object  
 2   area.code         5000 non-null   object  
 3   account.length   5000 non-null   int64  
 4   voice.plan       5000 non-null   object  
 5   voice.messages   5000 non-null   int64  
 6   intl.plan        5000 non-null   object  
 7   intl.mins        5000 non-null   float64 
 8   intl.calls       5000 non-null   int64  
 9   intl.charge      5000 non-null   float64 
 10  day.mins         5000 non-null   float64 
 11  day.calls        5000 non-null   int64  
 12  day.charge       4993 non-null   float64 
 13  eve.mins         4976 non-null   float64
```

```
14 eve.calls      5000 non-null  int64
15 eve.charge    5000 non-null  float64
16 night.mins    5000 non-null  float64
17 night.calls   5000 non-null  int64
18 night.charge  5000 non-null  float64
19 customer.calls 5000 non-null  int64
20 churn         5000 non-null  object
dtypes: float64(8), int64(8), object(5)
memory usage: 820.4+ KB
```

In [4]:

```
data.isna().sum()
```

Out[4]:

```
Unnamed: 0      0
state          0
area.code      0
account.length 0
voice.plan     0
voice.messages 0
intl.plan      0
intl.mins      0
intl.calls     0
intl.charge    0
day.mins       0
day.calls      0
day.charge     7
eve.mins       24
eve.calls      0
eve.charge     0
night.mins     0
night.calls    0
night.charge   0
customer.calls 0
churn          0
dtype: int64
```

In [5]:

```
data.dropna(inplace=True)
```

In [6]:

```
data.columns
```

Out[6]:

```
Index(['Unnamed: 0', 'state', 'area.code', 'account.length', 'voice.plan',
       'voice.messages', 'intl.plan', 'intl.mins', 'intl.calls', 'intl.charge',
       'day.mins', 'day.calls', 'day.charge', 'eve.mins', 'eve.calls',
       'eve.charge', 'night.mins', 'night.calls', 'night.charge',
       'customer.calls', 'churn'],
      dtype='object')
```

In [7]:

```
data.rename(columns={'area.code':'area_code','account.length':'account_length','voice.pl
                     'voice.messages':'voice_messages','intl.plan':'intl_plan','intl.min
                     'intl.calls':'intl_calls','intl.charge':'intl_charge','day.mins':'da
                     'day.charge':'day_charge','eve.mins':'eve_mins','eve.calls':'eve_cal
                     'night.mins':'night_mins','night.calls':'night_calls','night.charge'
                     },inplace=True)
```

```
data
```

Out[7]:

	Unnamed: 0	state	area_code	account_length	voice_plan	voice_messages	intl_plan	intl_min
0	1	KS	area_code_415	128	yes	25	no	10.
1	2	OH	area_code_415	107	yes	26	no	13.
2	3	NJ	area_code_415	137	no	0	no	12.
3	4	OH	area_code_408	84	no	0	yes	6.
4	5	OK	area_code_415	75	no	0	yes	10.
...
4995	4996	HI	area_code_408	50	yes	40	no	9.
4996	4997	WV	area_code_415	152	no	0	no	14.
4997	4998	DC	area_code_415	61	no	0	no	13.
4998	4999	DC	area_code_510	109	no	0	no	8.
4999	5000	VT	area_code_415	86	yes	34	no	9.

4969 rows × 21 columns

In [8]:

```
data.drop('Unnamed: 0',axis=1,inplace=True)  
data
```

Out[8]:

	state	area_code	account_length	voice_plan	voice_messages	intl_plan	intl_mins	intl_calls
0	KS	area_code_415	128	yes	25	no	10.0	3
1	OH	area_code_415	107	yes	26	no	13.7	3
2	NJ	area_code_415	137	no	0	no	12.2	5
3	OH	area_code_408	84	no	0	yes	6.6	7
4	OK	area_code_415	75	no	0	yes	10.1	3
...
4995	HI	area_code_408	50	yes	40	no	9.9	5
4996	WV	area_code_415	152	no	0	no	14.7	2
4997	DC	area_code_415	61	no	0	no	13.6	4
4998	DC	area_code_510	109	no	0	no	8.5	6
4999	VT	area_code_415	86	yes	34	no	9.3	16

4969 rows × 20 columns

In [9]:

```
le = preprocessing.LabelEncoder()  
objlist = ['area_code','intl_plan','voice_plan','churn','state']  
data[objlist] = data[objlist].apply(le.fit_transform)
```

```
In [10]:
```

```
data
```

```
Out[10]:
```

	state	area_code	account_length	voice_plan	voice_messages	intl_plan	intl_mins	intl_calls	in:
0	16	1	128	1	25	0	10.0	3	
1	35	1	107	1	26	0	13.7	3	
2	31	1	137	0	0	0	12.2	5	
3	35	0	84	0	0	1	6.6	7	
4	36	1	75	0	0	1	10.1	3	
...
4995	11	0	50	1	40	0	9.9	5	
4996	49	1	152	0	0	0	14.7	2	
4997	7	1	61	0	0	0	13.6	4	
4998	7	2	109	0	0	0	8.5	6	
4999	46	1	86	1	34	0	9.3	16	

4969 rows × 20 columns

```
In [11]:
```

```
x = data.iloc[:, 0:-1]
```

```
x
```

```
Out[11]:
```

	state	area_code	account_length	voice_plan	voice_messages	intl_plan	intl_mins	intl_calls	in:
0	16	1	128	1	25	0	10.0	3	
1	35	1	107	1	26	0	13.7	3	
2	31	1	137	0	0	0	12.2	5	
3	35	0	84	0	0	1	6.6	7	
4	36	1	75	0	0	1	10.1	3	
...
4995	11	0	50	1	40	0	9.9	5	
4996	49	1	152	0	0	0	14.7	2	
4997	7	1	61	0	0	0	13.6	4	
4998	7	2	109	0	0	0	8.5	6	
4999	46	1	86	1	34	0	9.3	16	

4969 rows × 19 columns

```
In [12]:
```

```
y = data.iloc[:, -1]
```

```
y
```

```
Out[12]:  
0      0  
1      0  
2      0  
3      0  
4      0  
..  
4995   0  
4996   1  
4997   0  
4998   0  
4999   0  
Name: churn, Length: 4969, dtype: int32
```

```
In [13]:
```

```
from sklearn.model_selection import train_test_split, cross_val_score
```

```
In [14]:
```

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3, random_state=15)
```

```
In [15]:
```

```
#Logistic Regression  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import classification_report  
from sklearn import preprocessing  
from sklearn import metrics
```

```
In [49]:
```

```
model=LogisticRegression()  
model.fit(x_train,y_train)  
y_pred=model.predict(x_test)  
print(classification_report(y_test,y_pred))  
from sklearn.metrics import precision_recall_fscore_support as score  
precision,recall,fscore,support = score(y_test,y_pred,average='macro')  
LR=fscore  
LR
```

	precision	recall	f1-score	support
0	0.87	0.99	0.93	1290
1	0.44	0.07	0.12	201
accuracy			0.86	1491
macro avg	0.65	0.53	0.52	1491
weighted avg	0.81	0.86	0.82	1491

```
Out[49]:
```

```
0.5227995509877178
```

```
In [17]:
```

```
from sklearn.metrics import confusion_matrix  
from sklearn.metrics import accuracy_score as ac  
confusion_matrix=confusion_matrix(y_test,y_pred)  
print(confusion_matrix)
```

```
[[1272  21]  
 [ 173  25]]
```

```
In [19]:
```

```

#KNN
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
%matplotlib inline
k_range = [2*i+1 for i in range(0,20)]
k_scores = []
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn,x_train , y_train, cv = 10)
    k_scores.append(scores.mean())
    print("K value=",k)
    model = KNeighborsClassifier(n_neighbors=k)
    model.fit(x_train,y_train)
    pred = model.predict(x_test)
    print(classification_report(y_test,pred))

```

K value= 1

	precision	recall	f1-score	support
0	0.90	0.88	0.89	1293
1	0.34	0.38	0.36	198
accuracy			0.82	1491
macro avg	0.62	0.63	0.63	1491
weighted avg	0.83	0.82	0.82	1491

K value= 3

	precision	recall	f1-score	support
0	0.90	0.96	0.93	1293
1	0.55	0.31	0.40	198
accuracy			0.88	1491
macro avg	0.73	0.64	0.66	1491
weighted avg	0.85	0.88	0.86	1491

K value= 5

	precision	recall	f1-score	support
0	0.90	0.98	0.94	1293
1	0.69	0.31	0.43	198
accuracy			0.89	1491
macro avg	0.80	0.65	0.68	1491
weighted avg	0.87	0.89	0.87	1491

K value= 7

	precision	recall	f1-score	support
0	0.90	0.98	0.94	1293
1	0.72	0.30	0.42	198
accuracy			0.89	1491
macro avg	0.81	0.64	0.68	1491
weighted avg	0.88	0.89	0.87	1491

K value= 9

	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293
1	0.76	0.28	0.41	198
accuracy			0.89	1491
macro avg	0.83	0.63	0.68	1491
weighted avg	0.88	0.89	0.87	1491

K value= 11

	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293
1	0.77	0.28	0.41	198
accuracy			0.89	1491
macro avg	0.83	0.63	0.68	1491
weighted avg	0.88	0.89	0.87	1491

K value= 13

	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293
1	0.77	0.28	0.41	198
accuracy			0.89	1491
macro avg	0.84	0.63	0.68	1491
weighted avg	0.88	0.89	0.87	1491

K value= 15

	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293
1	0.79	0.26	0.39	198
accuracy			0.89	1491
macro avg	0.84	0.63	0.67	1491
weighted avg	0.88	0.89	0.87	1491

K value= 17

	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293
1	0.80	0.26	0.39	198
accuracy			0.89	1491
macro avg	0.85	0.62	0.67	1491
weighted avg	0.88	0.89	0.87	1491

K value= 19

	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293
1	0.83	0.25	0.39	198
accuracy			0.89	1491
macro avg	0.86	0.62	0.66	1491
weighted avg	0.89	0.89	0.87	1491

K value= 21

	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293
1	0.83	0.25	0.39	198
accuracy			0.89	1491
macro avg	0.86	0.62	0.66	1491
weighted avg	0.89	0.89	0.87	1491

K value= 23

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.82	0.24	0.37	198
accuracy			0.89	1491
macro avg	0.86	0.61	0.65	1491
weighted avg	0.89	0.89	0.86	1491

K value= 25

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.80	0.23	0.35	198
accuracy			0.89	1491
macro avg	0.85	0.61	0.65	1491
weighted avg	0.88	0.89	0.86	1491

K value= 27

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.84	0.23	0.36	198
accuracy			0.89	1491
macro avg	0.87	0.61	0.65	1491
weighted avg	0.89	0.89	0.86	1491

K value= 29

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.83	0.22	0.35	198
accuracy			0.89	1491
macro avg	0.86	0.61	0.65	1491
weighted avg	0.88	0.89	0.86	1491

K value= 31

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.83	0.22	0.34	198
accuracy			0.89	1491
macro avg	0.86	0.61	0.64	1491

weighted avg 0.88 0.89 0.86 1491

K value= 33

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.84	0.21	0.34	198
accuracy			0.89	1491
macro avg	0.87	0.60	0.64	1491
weighted avg	0.88	0.89	0.86	1491

K value= 35

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.84	0.21	0.34	198
accuracy			0.89	1491
macro avg	0.87	0.60	0.64	1491
weighted avg	0.88	0.89	0.86	1491

K value= 37

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.84	0.21	0.34	198
accuracy			0.89	1491
macro avg	0.87	0.60	0.64	1491
weighted avg	0.88	0.89	0.86	1491


K value= 39

	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293
1	0.83	0.22	0.34	198
accuracy			0.89	1491
macro avg	0.86	0.61	0.64	1491
weighted avg	0.88	0.89	0.86	1491

In [20]:

```
plt.bar(k_range, k_scores)
plt.plot(k_range, k_scores, color = "red")

plt.xlabel('value of K for KNN')
plt.ylabel('Cross-Validation Accuracy')
plt.xticks(k_range)
plt.show()
```



In [21]:

```
model = KNeighborsClassifier(n_neighbors=5)
model.fit(x_train,y_train)
pred = model.predict(x_test)
```

In [51]:

```
print(classification_report(y_test,pred))
precision,recall,fscore,support = score(y_test,pred,average='macro')
KNN=fscore
KNN
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1290
1	0.93	0.78	0.85	201
accuracy			0.96	1491
macro avg	0.95	0.89	0.91	1491
weighted avg	0.96	0.96	0.96	1491

Out[51]:

0.9149499531734697

In [52]:

```
#Decision Tree
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
model = DecisionTreeClassifier(criterion='entropy',min_samples_split=5)
model.fit(x_train,y_train)
pred = model.predict(x_test)
metrics.accuracy_score(pred,y_test)
print(classification_report(y_test,pred))
precision,recall,fscore,support = score(y_test,pred,average='macro')
DT=fscore
DT
```

	precision	recall	f1-score	support
0	0.96	0.95	0.95	1290
1	0.69	0.73	0.71	201
accuracy			0.92	1491
macro avg	0.83	0.84	0.83	1491

```
weighted avg      0.92      0.92      0.92      1491
```

```
Out[52]:  
0.8327714404470512
```

```
In [53]:
```

```
# Random Forest Classification
```

```
from sklearn.model_selection import KFold  
from sklearn.model_selection import cross_val_score  
from sklearn.ensemble import RandomForestClassifier  
  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=7)  
  
num_trees = 140  
max_features = 6  
  
model = RandomForestClassifier(n_estimators=num_trees,max_samples=0.8, max_features=max_  
model.fit(x_train,y_train)  
  
pred = model.predict(x_test)  
  
print(classification_report(y_test,pred))  
precision,recall,fscore,support = score(y_test,pred,average='macro')  
RD=fscore  
RD
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1074
1	0.94	0.80	0.86	169
accuracy			0.97	1243
macro avg	0.95	0.90	0.92	1243
weighted avg	0.96	0.97	0.96	1243

```
Out[53]:  
0.9214157486080257
```

```
In [25]:
```

```
#SVM  
from sklearn import svm  
from sklearn.svm import SVC  
from sklearn.model_selection import GridSearchCV  
  
# clf = SVC()  
# param_grid = [{"kernel":["rbf"],'gamma':[50,5,10,0.5,1,0.001,0.0001,0.00001],'C':[1,15  
# gsv = GridSearchCV(clf,param_grid,cv=10)  
# gsv.fit(x_train,y_train)
```

```
In [55]:
```

```
clf = SVC(C= 15, gamma = 1,kernel="rbf")  
clf.fit(x_train,y_train)  
y_pred = clf.predict(x_test)  
print(classification_report(y_test, y_pred))  
precision,recall,fscore,support = score(y_test,y_pred,average='macro')  
SVM=fscore  
SVM
```

	precision	recall	f1-score	support
0	0.86	1.00	0.93	1074
1	0.00	0.00	0.00	169
accuracy			0.86	1243
macro avg	0.43	0.50	0.46	1243
weighted avg	0.75	0.86	0.80	1243

Out[55]:
0.46353042727665084

In [58]:

```
#ANN
import tensorflow as tf
model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(15, input_dim=19, activation = 'sigmoid'))
model.add(tf.keras.layers.Dense(10, activation='sigmoid'))
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))
```

In [59]:

```
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
dense_3 (Dense)	(None, 15)	300
dense_4 (Dense)	(None, 10)	160
dense_5 (Dense)	(None, 1)	11
<hr/>		
Total params: 471		
Trainable params: 471		
Non-trainable params: 0		

In [60]:

```
# Compile model
model.compile(loss ='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

In [61]:

```
history = model.fit(x,y,validation_split=0.20, epochs=50, batch_size=100)
```

```
Epoch 1/50
40/40 [=====] - 1s 7ms/step - loss: 1.0661 - accuracy: 0.1419 -
val_loss: 0.8822 - val_accuracy: 0.1419
Epoch 2/50
40/40 [=====] - 0s 5ms/step - loss: 0.7453 - accuracy: 0.3603 -
val_loss: 0.6273 - val_accuracy: 0.8561
Epoch 3/50
40/40 [=====] - 0s 4ms/step - loss: 0.5636 - accuracy: 0.8571 -
val_loss: 0.5116 - val_accuracy: 0.8581
Epoch 4/50
40/40 [=====] - 0s 4ms/step - loss: 0.4833 - accuracy: 0.8581 -
val_loss: 0.4599 - val_accuracy: 0.8581
Epoch 5/50
40/40 [=====] - 0s 4ms/step - loss: 0.4459 - accuracy: 0.8581 -
```

```
val_loss: 0.4344 - val_accuracy: 0.8581
Epoch 6/50
40/40 [=====] - 0s 6ms/step - loss: 0.4276 - accuracy: 0.8581 -
val_loss: 0.4217 - val_accuracy: 0.8581
Epoch 7/50
40/40 [=====] - 0s 6ms/step - loss: 0.4182 - accuracy: 0.8581 -
val_loss: 0.4151 - val_accuracy: 0.8581
Epoch 8/50
40/40 [=====] - 0s 5ms/step - loss: 0.4132 - accuracy: 0.8581 -
val_loss: 0.4115 - val_accuracy: 0.8581
Epoch 9/50
40/40 [=====] - 0s 6ms/step - loss: 0.4106 - accuracy: 0.8581 -
val_loss: 0.4095 - val_accuracy: 0.8581
Epoch 10/50
40/40 [=====] - 0s 4ms/step - loss: 0.4091 - accuracy: 0.8581 -
val_loss: 0.4086 - val_accuracy: 0.8581
Epoch 11/50
40/40 [=====] - 0s 4ms/step - loss: 0.4081 - accuracy: 0.8581 -
val_loss: 0.4079 - val_accuracy: 0.8581
Epoch 12/50
40/40 [=====] - 0s 6ms/step - loss: 0.4074 - accuracy: 0.8581 -
val_loss: 0.4075 - val_accuracy: 0.8581
Epoch 13/50
40/40 [=====] - 0s 6ms/step - loss: 0.4069 - accuracy: 0.8581 -
val_loss: 0.4071 - val_accuracy: 0.8581
Epoch 14/50
40/40 [=====] - 0s 5ms/step - loss: 0.4067 - accuracy: 0.8581 -
val_loss: 0.4069 - val_accuracy: 0.8581
Epoch 15/50
40/40 [=====] - 0s 4ms/step - loss: 0.4066 - accuracy: 0.8581 -
val_loss: 0.4063 - val_accuracy: 0.8581
Epoch 16/50
40/40 [=====] - 0s 4ms/step - loss: 0.4063 - accuracy: 0.8581 -
val_loss: 0.4064 - val_accuracy: 0.8581
Epoch 17/50
40/40 [=====] - 0s 4ms/step - loss: 0.4062 - accuracy: 0.8581 -
val_loss: 0.4063 - val_accuracy: 0.8581
Epoch 18/50
40/40 [=====] - 0s 4ms/step - loss: 0.4060 - accuracy: 0.8581 -
val_loss: 0.4056 - val_accuracy: 0.8581
Epoch 19/50
40/40 [=====] - 0s 4ms/step - loss: 0.4057 - accuracy: 0.8581 -
val_loss: 0.4060 - val_accuracy: 0.8581
Epoch 20/50
40/40 [=====] - 0s 4ms/step - loss: 0.4055 - accuracy: 0.8581 -
val_loss: 0.4059 - val_accuracy: 0.8581
Epoch 21/50
40/40 [=====] - 0s 4ms/step - loss: 0.4052 - accuracy: 0.8581 -
val_loss: 0.4045 - val_accuracy: 0.8581
Epoch 22/50
40/40 [=====] - 0s 4ms/step - loss: 0.4049 - accuracy: 0.8581 -
val_loss: 0.4041 - val_accuracy: 0.8581
Epoch 23/50
40/40 [=====] - 0s 4ms/step - loss: 0.4046 - accuracy: 0.8581 -
val_loss: 0.4040 - val_accuracy: 0.8581
Epoch 24/50
40/40 [=====] - 0s 4ms/step - loss: 0.4044 - accuracy: 0.8581 -
val_loss: 0.4038 - val_accuracy: 0.8581
Epoch 25/50
```

```
40/40 [=====] - 0s 5ms/step - loss: 0.4039 - accuracy: 0.8581 -  
val_loss: 0.4031 - val_accuracy: 0.8581  
Epoch 26/50  
40/40 [=====] - 0s 4ms/step - loss: 0.4035 - accuracy: 0.8581 -  
val_loss: 0.4032 - val_accuracy: 0.8581  
Epoch 27/50  
40/40 [=====] - 0s 3ms/step - loss: 0.4033 - accuracy: 0.8581 -  
val_loss: 0.4023 - val_accuracy: 0.8581  
Epoch 28/50  
40/40 [=====] - 0s 4ms/step - loss: 0.4026 - accuracy: 0.8581 -  
val_loss: 0.4017 - val_accuracy: 0.8581  
Epoch 29/50  
40/40 [=====] - 0s 4ms/step - loss: 0.4022 - accuracy: 0.8581 -  
val_loss: 0.4012 - val_accuracy: 0.8581  
Epoch 30/50  
40/40 [=====] - 0s 3ms/step - loss: 0.4018 - accuracy: 0.8581 -  
val_loss: 0.4012 - val_accuracy: 0.8581  
Epoch 31/50  
40/40 [=====] - 0s 3ms/step - loss: 0.4012 - accuracy: 0.8581 -  
val_loss: 0.4010 - val_accuracy: 0.8581  
Epoch 32/50  
40/40 [=====] - 0s 4ms/step - loss: 0.4005 - accuracy: 0.8581 -  
val_loss: 0.3983 - val_accuracy: 0.8581  
Epoch 33/50  
40/40 [=====] - 0s 3ms/step - loss: 0.3987 - accuracy: 0.8581 -  
val_loss: 0.3979 - val_accuracy: 0.8581  
Epoch 34/50  
40/40 [=====] - 0s 3ms/step - loss: 0.3976 - accuracy: 0.8581 -  
val_loss: 0.3955 - val_accuracy: 0.8581  
Epoch 35/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3968 - accuracy: 0.8581 -  
val_loss: 0.3983 - val_accuracy: 0.8581  
Epoch 36/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3961 - accuracy: 0.8581 -  
val_loss: 0.3939 - val_accuracy: 0.8581  
Epoch 37/50  
40/40 [=====] - 0s 3ms/step - loss: 0.3952 - accuracy: 0.8581 -  
val_loss: 0.3926 - val_accuracy: 0.8581  
Epoch 38/50  
40/40 [=====] - 0s 3ms/step - loss: 0.3944 - accuracy: 0.8581 -  
val_loss: 0.3918 - val_accuracy: 0.8581  
Epoch 39/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3926 - accuracy: 0.8581 -  
val_loss: 0.3938 - val_accuracy: 0.8581  
Epoch 40/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3920 - accuracy: 0.8581 -  
val_loss: 0.3883 - val_accuracy: 0.8581  
Epoch 41/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3920 - accuracy: 0.8581 -  
val_loss: 0.3892 - val_accuracy: 0.8581  
Epoch 42/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3911 - accuracy: 0.8581 -  
val_loss: 0.3878 - val_accuracy: 0.8581  
Epoch 43/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3906 - accuracy: 0.8581 -  
val_loss: 0.3858 - val_accuracy: 0.8581  
Epoch 44/50  
40/40 [=====] - 0s 4ms/step - loss: 0.3893 - accuracy: 0.8581 -  
val_loss: 0.3885 - val_accuracy: 0.8581
```

```
Epoch 45/50
40/40 [=====] - 0s 4ms/step - loss: 0.3879 - accuracy: 0.8581 -
val_loss: 0.3838 - val_accuracy: 0.8581
Epoch 46/50
40/40 [=====] - 0s 4ms/step - loss: 0.3866 - accuracy: 0.8581 -
val_loss: 0.3835 - val_accuracy: 0.8581
Epoch 47/50
40/40 [=====] - 0s 4ms/step - loss: 0.3857 - accuracy: 0.8581 -
val_loss: 0.3850 - val_accuracy: 0.8581
Epoch 48/50
40/40 [=====] - 0s 4ms/step - loss: 0.3844 - accuracy: 0.8581 -
val_loss: 0.3855 - val_accuracy: 0.8581
Epoch 49/50
40/40 [=====] - 0s 4ms/step - loss: 0.3848 - accuracy: 0.8581 -
val_loss: 0.3816 - val_accuracy: 0.8581
Epoch 50/50
40/40 [=====] - 0s 4ms/step - loss: 0.3823 - accuracy: 0.8581 -
val_loss: 0.3810 - val_accuracy: 0.8581
```

In [62]:

```
model.save_weights("mywt.kmw")
```

In [71]:

```
#evaluate the model
scores = model.evaluate(x,y)
print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
ANN=model.metrics_names[1], scores[1]*100
ANN=ANN[1]
ANN
```

```
156/156 [=====] - 0s 1ms/step - loss: 0.3811 - accuracy: 0.8581
accuracy: 85.81%
Out[71]:
85.8120322227478
```

In [33]:

```
#Naive_Bayes
from sklearn.naive_bayes import GaussianNB
nb = GaussianNB()
nb.fit(x_train, y_train)
```

Out[33]:

```
▼ GaussianNB
GaussianNB()
```

In [34]:

```
y_pred = nb.predict(x_test)
```

In [35]:

```
from sklearn.metrics import confusion_matrix
conf_matrix1=confusion_matrix(y_test, y_pred)
conf_matrix1
```

Out[35]:

```
array([[976,  98],
       [ 74,  95]], dtype=int64)
```

In [36]:

```

import seaborn as sns
%matplotlib inline
fig, ax= plt.subplots(1,1, figsize=(4,4), dpi=100)


xlabels=['0-Less than 50K','1-Above 50K']
ylabels=['0-Less than 50K','1-Above 50K']

sns.heatmap(conf_matrix1, annot=True, cbar=None, cmap="Purples", xticklabels=xlabels, yticklabels=ylabels, linewidths=1, linecolor='black')
ax.set_xlabel('Predicted Class')
ax.set_ylabel('True Class')

```

Out[36]:

Text(20.722222222222214, 0.5, 'True Class')



In [72]:

```

from sklearn.metrics import classification_report
names=['0-Less than 50K','1-Above 50K']
print(classification_report(y_test, y_pred,target_names=names))
precision,recall,fscore,support = score(y_test,y_pred,average='macro')
NB=fscore
NB

```

	precision	recall	f1-score	support
0-Less than 50K	0.86	1.00	0.93	1074
1-Above 50K	0.00	0.00	0.00	169
accuracy			0.86	1243
macro avg	0.43	0.50	0.46	1243
weighted avg	0.75	0.86	0.80	1243

Out[72]:

0.46353042727665084

In [73]:

#Bagging

```
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=2)

num_trees = 100
model = BaggingClassifier(max_samples=0.8, n_estimators=num_trees, random_state=8)

model.fit(x_train,y_train)
pred = model.predict(x_test)
print(classification_report(y_test,pred))
precision,recall,fscore,support = score(y_test,pred,average='macro')
BG=fscore
BG
```

	precision	recall	f1-score	support
0	0.96	0.99	0.98	1290
1	0.91	0.76	0.83	201
accuracy			0.96	1491
macro avg	0.93	0.87	0.90	1491
weighted avg	0.96	0.96	0.96	1491

Out[73]:
0.9012623649683373

In [74]:

#AdaBoost Classifier

```
from sklearn.ensemble import AdaBoostClassifier
clf = AdaBoostClassifier(random_state=96,base_estimator=RandomForestClassifier(random_st
    n_estimators=100,learning_rate=0.01)
clf.fit(x_train,y_train)
pred = clf.predict(x_test)
print(classification_report(y_test,pred))
report = classification_report(y_test,pred)
precision,recall,fscore,support = score(y_test,pred,average='macro')
ADB=fscore
ADB
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1290
1	0.93	0.78	0.85	201
accuracy			0.96	1491
macro avg	0.95	0.89	0.91	1491
weighted avg	0.96	0.96	0.96	1491

Out[74]:
0.9149499531734697

```
In [45]:
```

```
# from sklearn.metrics import precision_recall_fscore_support as score
# precision,recall,fscore,support = score(y_test,pred,average='macro')
# ada=fscore
```

```
In [46]:
```

```
# ada
```

```
Out[46]:
```

```
0.9149499531734697
```

Accuracy Scores All Models

```
In [79]:
```

```
# initialize list elements
data = [['Logistic Regression',LR],['KNearest Nighbour',KNN],
        ['Decision Tree',DT],['Random Forest',RD],['Support Vector Machine',SVM],[Artif
        ['Navie Bais',NB],['Bagging',BG],['AdaBoosting',ADB]]
# Create the pandas DataFrame with column name is provided explicitly
df = pd.DataFrame(data, columns=['Algorithm Names','Accuracy'])
df
```

```
Out[79]:
```

	Algorithm Names	Accuracy
0	Logistic Regression	0.522800
1	KNearest Nighbour	0.914950
2	Decision Tree	0.832771
3	Random Forest	0.921416
4	Support Vector Machine	0.463530
5	Artificial Neural Network	85.812032
6	Navie Bais	0.463530
7	Bagging	0.901262
8	AdaBoosting	0.914950

```
In [ ]:
```